

NAME

statfs, fstatfs – get filesystem statistics

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <sys/vfs.h> /* or <sys/statfs.h> */
```

```
int statfs(const char *path, struct statfs *buf);
```

```
int fstatfs(int fd, struct statfs *buf);
```

Unless you need the *f_type* field, you should use the standard **statvfs(3)** interface instead.

DESCRIPTION

The **statfs()** system call returns information about a mounted filesystem. *path* is the pathname of any file within the mounted filesystem. *buf* is a pointer to a *statfs* structure defined approximately as follows:

```
struct statfs {
    __fsword_t f_type;      /* Type of filesystem (see below) */
    __fsword_t f_bsize;     /* Optimal transfer block size */
    fsblkcnt_t f_blocks;    /* Total data blocks in filesystem */
    fsblkcnt_t f_bfree;     /* Free blocks in filesystem */
    fsblkcnt_t f_bavail;    /* Free blocks available to
                           unprivileged user */
    fsfilcnt_t f_files;     /* Total inodes in filesystem */
    fsfilcnt_t f_ffree;     /* Free inodes in filesystem */
    fsid_t      f_fsid;     /* Filesystem ID */
    __fsword_t f_namelen;   /* Maximum length of filenames */
    __fsword_t f_frsize;    /* Fragment size (since Linux 2.6) */
    __fsword_t f_flags;     /* Mount flags of filesystem
                           (since Linux 2.6.36) */
    __fsword_t f_spare[xxx]; /* Padding bytes reserved for future use */
};
```

The following filesystem types may appear in *f_type*:

ADFS_SUPER_MAGIC	0xadf5	
AFFS_SUPER_MAGIC	0xadff	
AFS_SUPER_MAGIC	0x5346414f	
ANON_INODE_FS_MAGIC	0x09041934	/* Anonymous inode FS (for pseudofiles that have no name; e.g., epoll, signalfd, bpf) */
AUTOFS_SUPER_MAGIC	0x0187	
BDEVFS_MAGIC	0x62646576	
BEFS_SUPER_MAGIC	0x42465331	
BFS_MAGIC	0x1badface	
BINFMTFS_MAGIC	0x42494e4d	
BPF_FS_MAGIC	0xcafe4a11	
BTRFS_SUPER_MAGIC	0x9123683e	
BTRFS_TEST_MAGIC	0x73727279	
CGROUP_SUPER_MAGIC	0x27e0eb	/* Cgroup pseudo FS */
CGROUP2_SUPER_MAGIC	0x63677270	/* Cgroup v2 pseudo FS */
CIFS_MAGIC_NUMBER	0xff534d42	
CODA_SUPER_MAGIC	0x73757245	
COH_SUPER_MAGIC	0x012ff7b7	
CRAMFS_MAGIC	0x28cd3d45	
DEBUGFS_MAGIC	0x64626720	

```

DEVFS_SUPER_MAGIC      0x1373      /* Linux 2.6.17 and earlier */
DEVPTS_SUPER_MAGIC     0x1cd1
ECRYPTFS_SUPER_MAGIC   0xf15f
EFIVARFS_MAGIC         0xde5e81e4
EFS_SUPER_MAGIC        0x00414a53
EXT_SUPER_MAGIC        0x137d      /* Linux 2.0 and earlier */
EXT2_OLD_SUPER_MAGIC   0xef51
EXT2_SUPER_MAGIC       0xef53
EXT3_SUPER_MAGIC       0xef53
EXT4_SUPER_MAGIC       0xef53
F2FS_SUPER_MAGIC       0xf2f52010
FUSE_SUPER_MAGIC       0x65735546
FUTEXFS_SUPER_MAGIC    0xbad1dea  /* Unused */
HFS_SUPER_MAGIC        0x4244
HOSTFS_SUPER_MAGIC     0x00c0ffee
HPFS_SUPER_MAGIC       0xf995e849
HUGETLBFS_MAGIC        0x958458f6
ISOFS_SUPER_MAGIC      0x9660
JFFS2_SUPER_MAGIC      0x72b6
JFS_SUPER_MAGIC        0x3153464a
MINIX_SUPER_MAGIC      0x137f      /* original minix FS */
MINIX_SUPER_MAGIC2     0x138f      /* 30 char minix FS */
MINIX2_SUPER_MAGIC     0x2468      /* minix V2 FS */
MINIX2_SUPER_MAGIC2    0x2478      /* minix V2 FS, 30 char names */
MINIX3_SUPER_MAGIC     0x4d5a      /* minix V3 FS, 60 char names */
MQQUEUE_MAGIC          0x19800202 /* POSIX message queue FS */
MSDOS_SUPER_MAGIC      0x4d44
MTD_INODE_FS_MAGIC     0x11307854
NCP_SUPER_MAGIC        0x564c
NFS_SUPER_MAGIC        0x6969
NILFS_SUPER_MAGIC      0x3434
NSFS_MAGIC             0x6e736673
NTFS_SB_MAGIC          0x5346544e
OCFS2_SUPER_MAGIC      0x7461636f
OPENPROM_SUPER_MAGIC   0x9fa1
OVERLAYFS_SUPER_MAGIC  0x794c7630
PIPEFS_MAGIC           0x50495045
PROC_SUPER_MAGIC       0x9fa0      /* /proc FS */
PSTOREFS_MAGIC         0x6165676c
QNX4_SUPER_MAGIC       0x002f
QNX6_SUPER_MAGIC       0x68191122
RAMFS_MAGIC            0x858458f6
REISERFS_SUPER_MAGIC   0x52654973
ROMFS_MAGIC            0x7275
SECURITYFS_MAGIC       0x73636673
SELINUX_MAGIC          0xf97cff8c
SMACK_MAGIC            0x43415d53
SMB_SUPER_MAGIC        0x517b
SMB2_MAGIC_NUMBER      0xfe534d42
SOCKFS_MAGIC           0x534f434b
SQUASHFS_MAGIC         0x73717368
SYSFS_MAGIC            0x62656572
SYSV2_SUPER_MAGIC      0x012ff7b6
SYSV4_SUPER_MAGIC      0x012ff7b5

```

TMPFS_MAGIC	0x01021994
TRACEFS_MAGIC	0x74726163
UDF_SUPER_MAGIC	0x15013346
UFS_MAGIC	0x00011954
USBDEVICE_SUPER_MAGIC	0x9fa2
V9FS_MAGIC	0x01021997
VXFS_SUPER_MAGIC	0xa501fcf5
XENFS_SUPER_MAGIC	0xabba1974
XENIX_SUPER_MAGIC	0x012ff7b4
XFS_SUPER_MAGIC	0x58465342
_XIAFS_SUPER_MAGIC	0x012fd16d /* Linux 2.0 and earlier */

Most of these MAGIC constants are defined in `/usr/include/linux/magic.h`, and some are hardcoded in kernel sources.

The `f_flags` field is a bit mask indicating mount options for the filesystem. It contains zero or more of the following bits:

ST_MANDLOCK

Mandatory locking is permitted on the filesystem (see **fcntl(2)**).

ST_NOATIME

Do not update access times; see **mount(2)**.

ST_NODEV

Disallow access to device special files on this filesystem.

ST_NODIRATIME

Do not update directory access times; see **mount(2)**.

ST_NOEXEC

Execution of programs is disallowed on this filesystem.

ST_NOSUID

The set-user-ID and set-group-ID bits are ignored by **exec(3)** for executable files on this filesystem

ST_RDONLY

This filesystem is mounted read-only.

ST_RELATIME

Update atime relative to mtime/ctime; see **mount(2)**.

ST_SYNCHRONOUS

Writes are synched to the filesystem immediately (see the description of **O_SYNC** in **open(2)**).

ST_NOSYMFOLLOW (since Linux 5.10)

Symbolic links are not followed when resolving paths; see **mount(2)**.

Nobody knows what `f_fsid` is supposed to contain (but see below).

Fields that are undefined for a particular filesystem are set to 0.

fstatfs() returns the same information about an open file referenced by descriptor `fd`.

RETURN VALUE

On success, zero is returned. On error, `-1` is returned, and `errno` is set to indicate the error.

ERRORS

EACCES

(**statfs()**) Search permission is denied for a component of the path prefix of `path`. (See also **path_resolution(7)**.)

EBADF

(**fstatfs()**) `fd` is not a valid open file descriptor.

EFAULT

buf or *path* points to an invalid address.

EINTR

The call was interrupted by a signal; see **signal(7)**.

EIO

An I/O error occurred while reading from the filesystem.

ELOOP

(**statfs()**) Too many symbolic links were encountered in translating *path*.

ENAMETOOLONG

(**statfs()**) *path* is too long.

ENOENT

(**statfs()**) The file referred to by *path* does not exist.

ENOMEM

Insufficient kernel memory was available.

ENOSYS

The filesystem does not support this call.

ENOTDIR

(**statfs()**) A component of the path prefix of *path* is not a directory.

EOVERFLOW

Some values were too large to be represented in the returned struct.

VERSIONS**The *f_fsid* field**

Solaris, Irix, and POSIX have a system call **statvfs(2)** that returns a *struct statvfs* (defined in `<sys/statvfs.h>`) containing an *unsigned long f_fsid*. Linux, SunOS, HP-UX, 4.4BSD have a system call **statfs()** that returns a *struct statfs* (defined in `<sys/vfs.h>`) containing a *fsid_t f_fsid*, where *fsid_t* is defined as *struct { int val[2]; }*. The same holds for FreeBSD, except that it uses the include file `<sys/mount.h>`.

The general idea is that *f_fsid* contains some random stuff such that the pair (*f_fsid*, *ino*) uniquely determines a file. Some operating systems use (a variation on) the device number, or the device number combined with the filesystem type. Several operating systems restrict giving out the *f_fsid* field to the superuser only (and zero it for unprivileged users), because this field is used in the filehandle of the filesystem when NFS-exported, and giving it out is a security concern.

Under some operating systems, the *fsid* can be used as the second argument to the **sysfs(2)** system call.

STANDARDS

Linux.

HISTORY

The Linux **statfs()** was inspired by the 4.4BSD one (but they do not use the same structure).

The original Linux **statfs()** and **fstatfs()** system calls were not designed with extremely large file sizes in mind. Subsequently, Linux 2.6 added new **statfs64()** and **fstatfs64()** system calls that employ a new structure, *statfs64*. The new structure contains the same fields as the original *statfs* structure, but the sizes of various fields are increased, to accommodate large file sizes. The glibc **statfs()** and **fstatfs()** wrapper functions transparently deal with the kernel differences.

LSB has deprecated the library calls **statfs()** and **fstatfs()** and tells us to use **statvfs(3)** and **fstatvfs(3)** instead.

NOTES

The *__fsword_t* type used for various fields in the *statfs* structure definition is a glibc internal type, not intended for public use. This leaves the programmer in a bit of a conundrum when trying to copy or compare these fields to local variables in a program. Using *unsigned int* for such variables suffices on most systems.

Some systems have only `<sys/vfs.h>`, other systems also have `<sys/statfs.h>`, where the former includes the latter. So it seems including the former is the best choice.

BUGS

From Linux 2.6.38 up to and including Linux 3.1, **fstatfs()** failed with the error **ENOSYS** for file descriptors created by **pipe(2)**.

SEE ALSO

stat(2), **statvfs(3)**, **path_resolution(7)**