

NAME

poll, ppoll – wait for some event on a file descriptor

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <poll.h>

int poll(struct pollfd *fds, nfds_t nfd, int timeout);

#define _GNU_SOURCE /* See feature_test_macros(7) */
#include <poll.h>

int ppoll(struct pollfd *fds, nfds_t nfd,
          const struct timespec *_Nullable tmo_p,
          const sigset_t *_Nullable sigmask);
```

DESCRIPTION

poll() performs a similar task to **select(2)**: it waits for one of a set of file descriptors to become ready to perform I/O. The Linux-specific **epoll(7)** API performs a similar task, but offers features beyond those found in **poll()**.

The set of file descriptors to be monitored is specified in the *fds* argument, which is an array of structures of the following form:

```
struct pollfd {
    int    fd;           /* file descriptor */
    short  events;       /* requested events */
    short  revents;      /* returned events */
};
```

The caller should specify the number of items in the *fds* array in *nfd*.

The field *fd* contains a file descriptor for an open file. If this field is negative, then the corresponding *events* field is ignored and the *revents* field returns zero. (This provides an easy way of ignoring a file descriptor for a single **poll()** call: simply set the *fd* field to its bitwise complement.)

The field *events* is an input parameter, a bit mask specifying the events the application is interested in for the file descriptor *fd*. This field may be specified as zero, in which case the only events that can be returned in *revents* are **POLLHUP**, **POLLERR**, and **POLLNVAL** (see below).

The field *revents* is an output parameter, filled by the kernel with the events that actually occurred. The bits returned in *revents* can include any of those specified in *events*, or one of the values **POLLERR**, **POLLHUP**, or **POLLNVAL**. (These three bits are meaningless in the *events* field, and will be set in the *revents* field whenever the corresponding condition is true.)

If none of the events requested (and no error) has occurred for any of the file descriptors, then **poll()** blocks until one of the events occurs.

The *timeout* argument specifies the number of milliseconds that **poll()** should block waiting for a file descriptor to become ready. The call will block until either:

- a file descriptor becomes ready;
- the call is interrupted by a signal handler; or
- the timeout expires.

Being "ready" means that the requested operation will not block; thus, **poll()**ing regular files, block devices, and other files with no reasonable polling semantic *always* returns instantly as ready to read and write.

Note that the *timeout* interval will be rounded up to the system clock granularity, and kernel scheduling delays mean that the blocking interval may overrun by a small amount. Specifying a negative value in *timeout* means an infinite timeout. Specifying a *timeout* of zero causes **poll()** to return immediately, even if no file descriptors are ready.

The bits that may be set/returned in *events* and *revents* are defined in `<poll.h>`:

POLLIN

There is data to read.

POLLPRI

There is some exceptional condition on the file descriptor. Possibilities include:

- There is out-of-band data on a TCP socket (see **tcp(7)**).
- A pseudoterminal master in packet mode has seen a state change on the slave (see **ioctl_tty(2)**).
- A *cgroup.events* file has been modified (see **cgroups(7)**).

POLLOUT

Writing is now possible, though a write larger than the available space in a socket or pipe will still block (unless **O_NONBLOCK** is set).

POLLRDHUP (since Linux 2.6.17)

Stream socket peer closed connection, or shut down writing half of connection. The **_GNU_SOURCE** feature test macro must be defined (before including *any* header files) in order to obtain this definition.

POLLERR

Error condition (only returned in *revents*; ignored in *events*). This bit is also set for a file descriptor referring to the write end of a pipe when the read end has been closed.

POLLHUP

Hang up (only returned in *revents*; ignored in *events*). Note that when reading from a channel such as a pipe or a stream socket, this event merely indicates that the peer closed its end of the channel. Subsequent reads from the channel will return 0 (end of file) only after all outstanding data in the channel has been consumed.

POLLNVAL

Invalid request: *fd* not open (only returned in *r events*; ignored in *events*).

When compiling with **_XOPEN_SOURCE** defined, one also has the following, which convey no further information beyond the bits listed above:

POLLRDNORM

Equivalent to **POLLIN**.

POLLRDBAND

Priority band data can be read (generally unused on Linux).

POLLWRNORM

Equivalent to **POLLOUT**.

POLLWRBAND

Priority data may be written.

Linux also knows about, but does not use **POLLMSG**.

ppoll()

The relationship between **poll()** and **ppoll()** is analogous to the relationship between **select(2)** and **pselect(2)**: like **pselect(2)**, **ppoll()** allows an application to safely wait until either a file descriptor becomes ready or until a signal is caught.

Other than the difference in the precision of the *timeout* argument, the following **ppoll()** call:

```
ready = ppoll(&fds, nfd, tmo_p, &sigmask);
```

is nearly equivalent to *atomically* executing the following calls:

```
sigset_t origmask;
int timeout;
```

```

timeout = (tmo_p == NULL) ? -1 :
           (tmo_p->tv_sec * 1000 + tmo_p->tv_nsec / 1000000);
pthread_sigmask(SIG_SETMASK, &sigmask, &origmask);
ready = poll(&fds, nfds, timeout);
pthread_sigmask(SIG_SETMASK, &origmask, NULL);

```

The above code segment is described as *nearly* equivalent because whereas a negative *timeout* value for **poll()** is interpreted as an infinite timeout, a negative value expressed in **tmo_p* results in an error from **ppoll()**.

See the description of **pselect(2)** for an explanation of why **ppoll()** is necessary.

If the *sigmask* argument is specified as NULL, then no signal mask manipulation is performed (and thus **ppoll()** differs from **poll()** only in the precision of the *timeout* argument).

The *tmo_p* argument specifies an upper limit on the amount of time that **ppoll()** will block. This argument is a pointer to a **timespec(3)** structure.

If *tmo_p* is specified as NULL, then **ppoll()** can block indefinitely.

RETURN VALUE

On success, **poll()** returns a nonnegative value which is the number of elements in the *pollfds* whose *revents* fields have been set to a nonzero value (indicating an event or an error). A return value of zero indicates that the system call timed out before any file descriptors became ready.

On error, -1 is returned, and *errno* is set to indicate the error.

ERRORS

EFAULT

fds points outside the process's accessible address space. The array given as argument was not contained in the calling program's address space.

EINTR

A signal occurred before any requested event; see **signal(7)**.

EINVAL

The *nfds* value exceeds the **RLIMIT_NOFILE** value.

EINVAL

(**ppoll()**) The timeout value expressed in **tmo_p* is invalid (negative).

ENOMEM

Unable to allocate memory for kernel data structures.

VERSIONS

On some other UNIX systems, **poll()** can fail with the error **EAGAIN** if the system fails to allocate kernel-internal resources, rather than **ENOMEM** as Linux does. POSIX permits this behavior. Portable programs may wish to check for **EAGAIN** and loop, just as with **EINTR**.

Some implementations define the nonstandard constant **INFTIM** with the value -1 for use as a *timeout* for **poll()**. This constant is not provided in glibc.

C library/kernel differences

The Linux **ppoll()** system call modifies its *tmo_p* argument. However, the glibc wrapper function hides this behavior by using a local variable for the timeout argument that is passed to the system call. Thus, the glibc **ppoll()** function does not modify its *tmo_p* argument.

The raw **ppoll()** system call has a fifth argument, *size_t sigsetsize*, which specifies the size in bytes of the *sigmask* argument. The glibc **ppoll()** wrapper function specifies this argument as a fixed value (equal to `sizeof(kernel_sigset_t)`). See **sigprocmask(2)** for a discussion on the differences between the kernel and the libc notion of the sigset.

STANDARDS**poll()** POSIX.1-2008.**ppoll()** Linux.**HISTORY****poll()** POSIX.1-2001. Linux 2.1.23.

On older kernels that lack this system call, the glibc **poll()** wrapper function provides emulation using **select(2)**.

ppoll() Linux 2.6.16, glibc 2.4.**NOTES**

The operation of **poll()** and **ppoll()** is not affected by the **O_NONBLOCK** flag.

For a discussion of what may happen if a file descriptor being monitored by **poll()** is closed in another thread, see **select(2)**.

BUGS

See the discussion of spurious readiness notifications under the **BUGS** section of **select(2)**.

EXAMPLES

The program below opens each of the files named in its command-line arguments and monitors the resulting file descriptors for readiness to read (**POLLIN**). The program loops, repeatedly using **poll()** to monitor the file descriptors, printing the number of ready file descriptors on return. For each ready file descriptor, the program:

- displays the returned *revents* field in a human-readable form;
- if the file descriptor is readable, reads some data from it, and displays that data on standard output; and
- if the file descriptor was not readable, but some other event occurred (presumably **POLLHUP**), closes the file descriptor.

Suppose we run the program in one terminal, asking it to open a FIFO:

```
$ mkfifo myfifo
$ ./poll_input myfifo
```

In a second terminal window, we then open the FIFO for writing, write some data to it, and close the FIFO:

```
$ echo aaaaabbbbbccccc > myfifo
```

In the terminal where we are running the program, we would then see:

```
Opened "myfifo" on fd 3
About to poll()
Ready: 1
  fd=3; events: POLLIN POLLHUP
  read 10 bytes: aaaaabbbbb
About to poll()
Ready: 1
  fd=3; events: POLLIN POLLHUP
  read 6 bytes: ccccc

About to poll()
Ready: 1
  fd=3; events: POLLHUP
  closing fd 3
All file descriptors closed; bye
```

In the above output, we see that **poll()** returned three times:

- On the first return, the bits returned in the *revents* field were **POLLIN**, indicating that the file descriptor is readable, and **POLLHUP**, indicating that the other end of the FIFO has been closed. The program then consumed some of the available input.
- The second return from **poll()** also indicated **POLLIN** and **POLLHUP**; the program then consumed the last of the available input.
- On the final return, **poll()** indicated only **POLLHUP** on the FIFO, at which point the file descriptor was closed and the program terminated.

Program source

```
/* poll_input.c

Licensed under GNU General Public License v2 or later.
*/
#include <fcntl.h>
#include <poll.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

#define errExit(msg)    do { perror(msg); exit(EXIT_FAILURE); \
                        } while (0)

int
main(int argc, char *argv[])
{
    int            ready;
    char          buf[10];
    nfds_t         num_open_fds, nfds;
    ssize_t        s;
    struct pollfd  *pfds;

    if (argc < 2) {
        fprintf(stderr, "Usage: %s file...\n", argv[0]);
        exit(EXIT_FAILURE);
    }

    num_open_fds = nfds = argc - 1;
    pfd = calloc(nfds, sizeof(struct pollfd));
    if (pfd == NULL)
        errExit("malloc");

    /* Open each file on command line, and add it to 'pfd' array. */

    for (nfds_t j = 0; j < nfds; j++) {
        pfd[j].fd = open(argv[j + 1], O_RDONLY);
        if (pfd[j].fd == -1)
            errExit("open");

        printf("Opened \"%s\" on fd %d\n", argv[j + 1], pfd[j].fd);

        pfd[j].events = POLLIN;
    }
}
```

```

/* Keep calling poll() as long as at least one file descriptor is
open. */

while (num_open_fds > 0) {
    printf("About to poll()\n");
    ready = poll(pfds, nfds, -1);
    if (ready == -1)
        errExit("poll");

    printf("Ready: %d\n", ready);

    /* Deal with array returned by poll(). */

    for (nfds_t j = 0; j < nfds; j++) {
        if (pfds[j].revents != 0) {
            printf("  fd=%d; events: %s%s%s\n", pfds[j].fd,
                (pfds[j].revents & POLLIN) ? "POLLIN " : "",
                (pfds[j].revents & POLLHUP) ? "POLLHUP " : "",
                (pfds[j].revents & POLLERR) ? "POLLERR " : "");

            if (pfds[j].revents & POLLIN) {
                s = read(pfds[j].fd, buf, sizeof(buf));
                if (s == -1)
                    errExit("read");
                printf("    read %zd bytes: %.*s\n",
                    s, (int) s, buf);
            } else {
                /* POLLERR | POLLHUP */
                printf("    closing fd %d\n", pfds[j].fd);
                if (close(pfds[j].fd) == -1)
                    errExit("close");
                num_open_fds--;
            }
        }
    }

    printf("All file descriptors closed; bye\n");
    exit(EXIT_SUCCESS);
}

```

SEE ALSO

restart_syscall(2), select(2), select_tut(2), timespec(3), epoll(7), time(7)