

NAME

sched_setscheduler, sched_getscheduler – set and get scheduling policy/parameters

LIBRARY

Standard C library (*libc*, *-lc*)

SYNOPSIS

```
#include <sched.h>

int sched_setscheduler(pid_t pid, int policy,
                      const struct sched_param *param);
int sched_getscheduler(pid_t pid);
```

DESCRIPTION

The **sched_setscheduler()** system call sets both the scheduling policy and parameters for the thread whose ID is specified in *pid*. If *pid* equals zero, the scheduling policy and parameters of the calling thread will be set.

The scheduling parameters are specified in the *param* argument, which is a pointer to a structure of the following form:

```
struct sched_param {
    ...
    int sched_priority;
    ...
};
```

In the current implementation, the structure contains only one field, *sched_priority*. The interpretation of *param* depends on the selected policy.

Currently, Linux supports the following "normal" (i.e., non-real-time) scheduling policies as values that may be specified in *policy*:

SCHED_OTHER

the standard round-robin time-sharing policy;

SCHED_BATCH

for "batch" style execution of processes; and

SCHED_IDLE for running very low priority background jobs.

For each of the above policies, *param->sched_priority* must be 0.

Various "real-time" policies are also supported, for special time-critical applications that need precise control over the way in which runnable threads are selected for execution. For the rules governing when a process may use these policies, see **sched(7)**. The real-time policies that may be specified in *policy* are:

SCHED_FIFO a first-in, first-out policy; and

SCHED_RR a round-robin policy.

For each of the above policies, *param->sched_priority* specifies a scheduling priority for the thread. This is a number in the range returned by calling **sched_get_priority_min(2)** and **sched_get_priority_max(2)** with the specified *policy*. On Linux, these system calls return, respectively, 1 and 99.

Since Linux 2.6.32, the **SCHED_RESET_ON_FORK** flag can be ORed in *policy* when calling **sched_setscheduler()**. As a result of including this flag, children created by **fork(2)** do not inherit privileged scheduling policies. See **sched(7)** for details.

sched_getscheduler() returns the current scheduling policy of the thread identified by *pid*. If *pid* equals zero, the policy of the calling thread will be retrieved.

RETURN VALUE

On success, **sched_setscheduler()** returns zero. On success, **sched_getscheduler()** returns the policy for the thread (a nonnegative integer). On error, both calls return *-1*, and *errno* is set to indicate the error.

ERRORS**EINVAL**

Invalid arguments: *pid* is negative or *param* is NULL.

EINVAL

(**sched_setscheduler**()) *policy* is not one of the recognized policies.

EINVAL

(**sched_setscheduler**()) *param* does not make sense for the specified *policy*.

EPERM

The calling thread does not have appropriate privileges.

ESRCH

The thread whose ID is *pid* could not be found.

VERSIONS

POSIX.1 does not detail the permissions that an unprivileged thread requires in order to call **sched_setscheduler**(), and details vary across systems. For example, the Solaris 7 manual page says that the real or effective user ID of the caller must match the real user ID or the save set-user-ID of the target.

The scheduling policy and parameters are in fact per-thread attributes on Linux. The value returned from a call to **gettid**(2) can be passed in the argument *pid*. Specifying *pid* as 0 will operate on the attributes of the calling thread, and passing the value returned from a call to **getpid**(2) will operate on the attributes of the main thread of the thread group. (If you are using the POSIX threads API, then use **pthread_setschedparam**(3), **pthread_getschedparam**(3), and **pthread_setschedprio**(3), instead of the **sched_***(2) system calls.)

STANDARDS

POSIX.1-2008 (but see BUGS below).

SCHED_BATCH and **SCHED_IDLE** are Linux-specific.

HISTORY

POSIX.1-2001.

NOTES

Further details of the semantics of all of the above "normal" and "real-time" scheduling policies can be found in the **sched**(7) manual page. That page also describes an additional policy, **SCHED_DEADLINE**, which is settable only via **sched_setattr**(2).

POSIX systems on which **sched_setscheduler**() and **sched_getscheduler**() are available define **_POSIX_PRIORITY_SCHEDULING** in *<unistd.h>*.

BUGS

POSIX.1 says that on success, **sched_setscheduler**() should return the previous scheduling policy. Linux **sched_setscheduler**() does not conform to this requirement, since it always returns 0 on success.

SEE ALSO

chrt(1), **nice**(2), **sched_get_priority_max**(2), **sched_get_priority_min**(2), **sched_getaffinity**(2), **sched_getattr**(2), **sched_getparam**(2), **sched_rr_get_interval**(2), **sched_setaffinity**(2), **sched_setattr**(2), **sched_setparam**(2), **sched_yield**(2), **setpriority**(2), **capabilities**(7), **cpuset**(7), **sched**(7)